

Image Processing I

Computer Vision: AI3604

Image Processing I

Transform image to new one that is easier to manipulate.

Topics:

(1) Pixel Processing

(2) Convolution

(3) Linear Filtering

(4) Non-Linear Filtering

(5) Correlation

} **Lecture 1**

Computer Vision: Algorithms and Applications (Chapter 3.2)
Szelinski, 2011 (available online)

Image Processing II

Transform image to new one that is easier to manipulate.

Topics:

(6) Frequency Representation of Signals

(7) Fourier Transform


(8) Convolution \leftrightarrow Fourier Transform

(9) Deconvolution in Frequency Domain

(10) Binary Image Processing

Lecture 2

Image as a Function



$f(x, y)$ is the image intensity at position (x, y)

Image Processing

Transformation t of one image f to another image g

$$g(x, y) = t(f(x, y))$$

Point (Pixel) Processing



Darken ($f - 128$)



Original (f)



Lighten ($f + 128$)



Invert ($255 - f$)

Point (Pixel) Processing



Low Contrast ($f/2$)



Original (f)



High Contrast ($f * 2$)



Gray
($0.3f_R + 0.6f_G + 0.1f_B$)

Linear Shift Invariant System







Study of Linear Shift Invariant Systems (**LSIS**) leads to useful image processing algorithms.


LSIS: Linearity



LSIS: Shift Invariance



Ideal Lens is an LSIS



Defocused Image (g) is a Processed version of Focused Image (f)


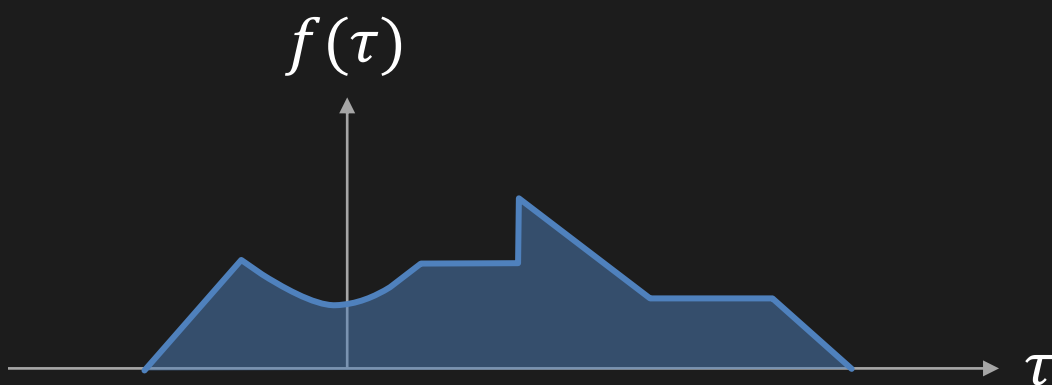
Linearity: Brightness variation

Shift invariance: Scene movement

Convolution (Important Concept)

Convolution of two functions $f(x)$ and $h(x)$



$$g(x) = f(x) * h(x) = \int_{-\infty}^{\infty} f(\tau)h(x - \tau) d\tau$$



Convolution (Important Concept)

Convolution of two functions $f(x)$ and $h(x)$



$$g(x) = f(x) * h(x) = \int_{-\infty}^{\infty} f(\tau)h(x - \tau) d\tau$$



Convolution (Important Concept)

Convolution of two functions $f(x)$ and $h(x)$


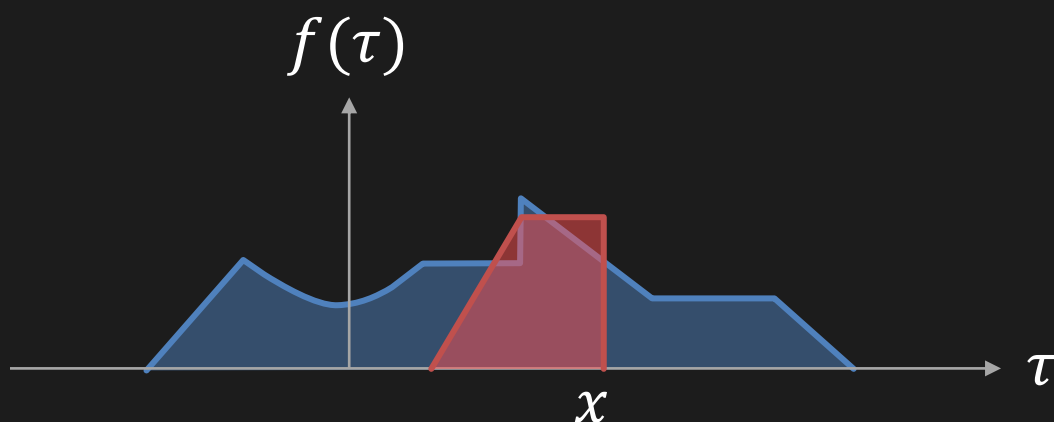
$$g(x) = f(x) * h(x) = \int_{-\infty}^{\infty} f(\tau)h(x - \tau) d\tau$$



Convolution (Important Concept)

Convolution of two functions $f(x)$ and $h(x)$


$$g(x) = f(x) * h(x) = \int_{-\infty}^{\infty} f(\tau)h(x - \tau) d\tau$$



Convolution (Important Concept)

Convolution of two functions $f(x)$ and $h(x)$


$$g(x) = f(x) * h(x) = \int_{-\infty}^{\infty} f(\tau)h(x - \tau) d\tau$$



Convolution (Important Concept)

Convolution of two functions $f(x)$ and $h(x)$



$$g(x) = f(x) * h(x) = \int_{-\infty}^{\infty} f(\tau)h(x - \tau) d\tau$$



Convolution (Important Concept)

Convolution of two functions $f(x)$ and $h(x)$


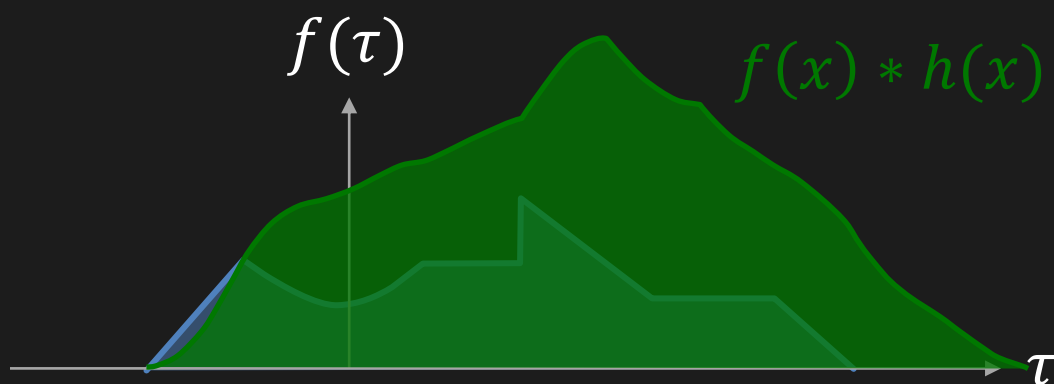
$$g(x) = f(x) * h(x) = \int_{-\infty}^{\infty} f(\tau)h(x - \tau) d\tau$$



Convolution (Important Concept)


Convolution of two functions $f(x)$ and $h(x)$

$$g(x) = f(x) * h(x) = \int_{-\infty}^{\infty} f(\tau)h(x - \tau) d\tau$$



Convolution implies LSIS and LSIS implies Convolution


Convolution: Example



$f(x)$




$h(x)$




$f(x) * h(x)$


Convolution: Example



$f(x)$



$h(x)$



$f(x) * h(x)$

Can we find h ?

$f \longrightarrow \boxed{h} \longrightarrow g$ $g(x) = \int_{-\infty}^{\infty} f(\tau)h(x - \tau) d\tau$

What input f will produce output $g = h$?


$$h(x) = \int_{-\infty}^{\infty} ?(\tau)h(x - \tau) d\tau$$

Unit Impulse Function

$$\delta(x) = \begin{cases} 1/2\varepsilon, & |x| \leq \varepsilon \\ 0, & |x| > \varepsilon \end{cases}$$

$\varepsilon \rightarrow 0$

$$\int_{-\infty}^{\infty} \delta(\tau) d\tau = \frac{1}{2\varepsilon} \cdot 2\varepsilon = 1$$



$$\int_{-\infty}^{\infty} \delta(\tau) b(x - \tau) d\tau = b(x)$$


Shifting Property

Impulse Response



$$g(x) = f(x) * h(x)$$


$$g(x) = \int_{-\infty}^{\infty} f(\tau) h(x - \tau) d\tau$$



$$h(x) = \delta(x) * h(x)$$

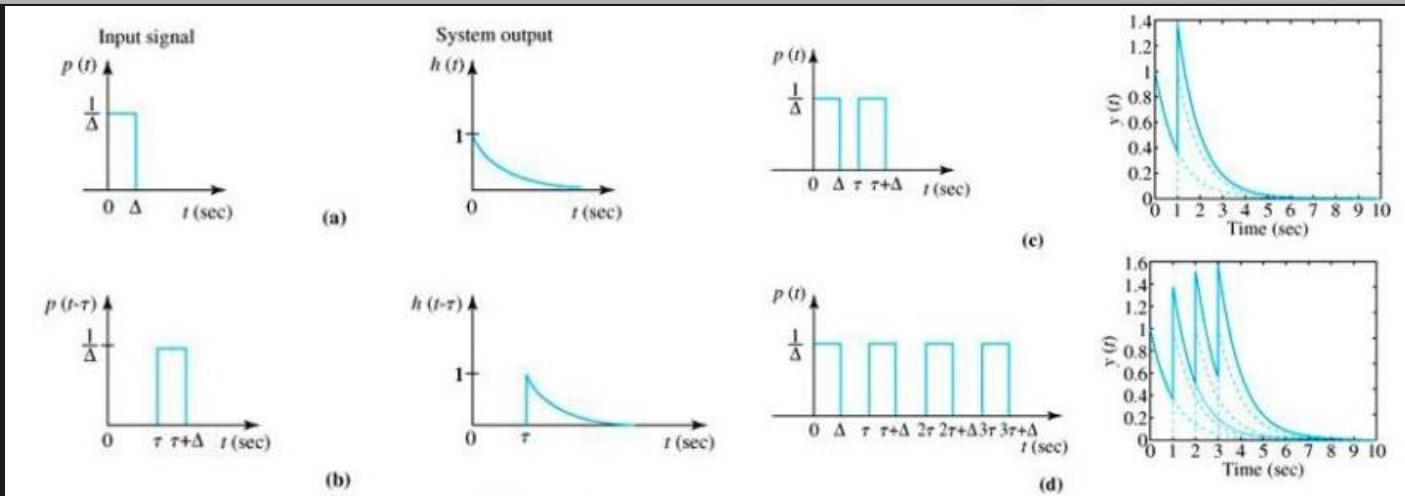
$$h(x) = \int_{-\infty}^{\infty} \delta(\tau) h(x - \tau) d\tau$$

Impulse Response of Human Eye

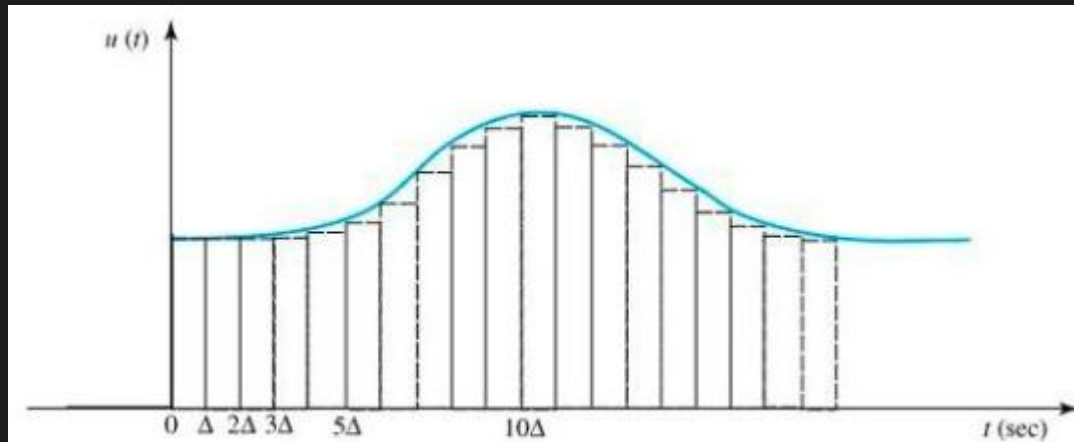


Human Eye PSF

LSIS implies Convolution



$$p(t) \rightarrow h(t), p(t - \Delta) \rightarrow h(t - \Delta), p(t) + p(t - \Delta) \rightarrow h(t) + h(t - \Delta)$$




$$\sum_{i=0}^{\infty} u(i\Delta)p(t-i\Delta) \rightarrow \sum_{i=0}^{\infty} u(i\Delta)h(t-i\Delta) \sim \int_0^{\infty} u(\tau)h(t-\tau)d\tau$$

Properties of Convolution

Commutative $a * b = b * a$


Associative $(a * b) * c = a * (b * c)$

Cascaded System



2D Convolution


LSIS:



Convolution:


$$g(x, y) = \int_{-\infty}^{\infty} \int_{-\infty}^{\infty} f(\tau, \mu) h(x - \tau, y - \mu) d\tau d\mu$$

Convolution with Discrete Images



$$g[i, j] = \sum_{m=1}^M \sum_{n=1}^N f[m, n] h[i - m, j - n]$$

Border Problem




Solution:

- Ignore Border
- Pad with Constant Value
- Pad with Reflection


Example: Impulse Filter

Input



$f(x, y)$

*



=

Output



$f(x, y)$


Example: Image Shift

Input




$$f(x, y)$$

*



=


Output



$$f(x - u, y - v)$$

Example: Averaging

Input



$f(x, y)$

$$* \begin{array}{c} \text{5x5 grid with a white center} \\ \text{"Box Filter"} \\ 5 \times 5 \end{array} =$$

Output



$g(x, y)$

Result Image is Saturated. Why?

Example: Averaging

Input




$f(x, y)$

$$* \begin{array}{c} \text{5 x 5} \\ \text{"Box Filter"} \end{array} =$$

$a(x, y)$


Output




$g(x, y)$

Sum of all the Filter (Kernel) Weights should be 1.



Linear filtering (no change)



original




—



Filtered
(no change)

(Swiped from Bill Freeman)


Linear filtering



original




—





?

(Swiped from Bill Freeman)

(remember blurring)




original




Blurred (filter applied in both dimensions).


Sharpening



original




—




Sharpened
original

borrowed from D. Kriegman

Sharpening



before



after

Smoothing With Box Filter

Input




$$f(x, y)$$

$$* \begin{array}{c} \square \\ \square \end{array} =$$

“Box Filter”
21 x 21

$$a(x, y)$$

Output



$$g(x, y)$$

Image smoothed with a box filter does not look “natural.”
Has blocky artifacts.


Smoothing With “Fuzzy” Filter

Input



$f(x, y)$


*



=

“Fuzzy Filter”
21 x 21

Output





$g(x, y)$

Gaussian Kernel: A Fuzzy Filter

$$n_{\sigma}[i, j] = \frac{1}{2\pi\sigma^2} e^{-\frac{1}{2}\left(\frac{i^2+j^2}{\sigma^2}\right)}$$

σ^2 : Variance



Rule of Thumb: Set Kernel Size $K \approx 2\pi\sigma$


Gaussian Smoothing

Input



$f(x, y)$

*



$\sigma = 4$

=

Output



$g(x, y)$

Larger the Kernel (or σ), More the Blurring


Gaussian Smoothing

Input



$f(x, y)$

*




$\sigma = 8$

$n_8(x, y)$

=

Output



$g(x, y)$

Larger the Kernel (or σ), More the Blurring


Gaussian Smoothing

Input



$f(x, y)$

*



$\sigma = 16$

$n_{16}(x, y)$

=

Output



$g(x, y)$

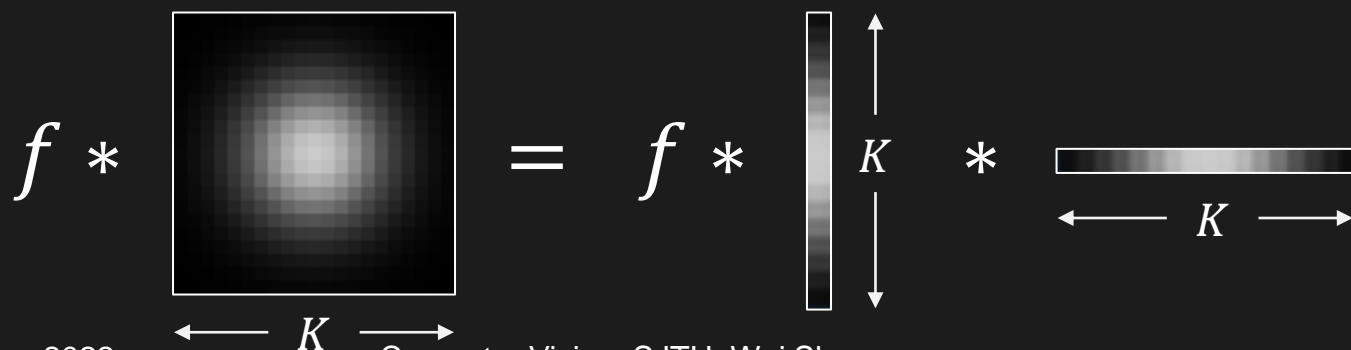
Larger the Kernel (or σ), More the Blurring

Gaussian Smoothing is Separable

$$g[i, j] = \frac{1}{2\pi\sigma^2} \sum_{m=-K/2}^{K/2} \sum_{n=-K/2}^{K/2} e^{-\frac{1}{2}\left(\frac{m^2+n^2}{\sigma^2}\right)} f[i-m, j-n]$$

$$g[i, j] = \frac{1}{2\pi\sigma^2} \sum_{m=-K/2}^{K/2} e^{-\frac{1}{2}\left(\frac{m^2}{\sigma^2}\right)} \cdot \sum_{n=-K/2}^{K/2} e^{-\frac{1}{2}\left(\frac{n^2}{\sigma^2}\right)} f[i-m, j-n]$$

Using One 2D Gaussian Filter \equiv Using Two 1D Gaussian Filters



Gaussian Smoothing is Separable

Using One 2D Gaussian Filter \equiv Using Two 1D Gaussian Filters

$$f * \begin{array}{c} \text{2D Gaussian Filter} \\ \leftarrow K \rightarrow \end{array} = f * \begin{array}{c} \text{1D Gaussian Filter} \\ \uparrow K \\ \downarrow K \end{array} * \begin{array}{c} \text{1D Gaussian Filter} \\ \leftarrow K \rightarrow \end{array}$$

Which one is faster? Why?


K^2 Multiplications

$K^2 - 1$ Additions

$2K$ Multiplications

$2(K - 1)$ Additions

Smoothing to Remove Image Noise



Problem with Smoothing:

- Sensitive to Outliers (Noise)
- Smoothens Edges (Blur)

Median Filtering

1. Sort the K^2 values in window centered at the pixel
2. Assign the Middle value (**Median**) to pixel





Image with
Salt and Pepper Noise



Median Filtered
Image ($K = 3$)


Non-linear Operation
(Cannot be implemented using Convolution)

Median Filtering

Not Effective when Image Noise is not a Simple Salt and Pepper Noise.



Image with Noise



Median Filtered
Image ($K = 7$)


Larger K causes Blurring of Image Detail

Median Filtering

Not Effective when Image Noise is not a Simple Salt and Pepper Noise.




Image with Noise



Median Filtered
Image ($K = 11$)


Larger K causes Blurring of Image Detail

Revisiting Gaussian Smoothing



Same Gaussian Kernel is used Everywhere
Blurs Across Edges


Blur Similar Pixels Only



“**Bias**” Gaussian Kernel such that pixels not similar in intensity to the center pixel receive a lower weight.


Bilateral Filter: Start With Gaussian

$$g[i, j] = \frac{1}{W_s} \sum_m \sum_n f[m, n] \underbrace{n_{\sigma_s}[i - m, j - n]}_{\text{Spatial Gaussian}}$$



Gaussian Blurs Across Edges

Bilateral Filter: Add Bias to Gaussian



Bilateral Filter: Summary

$$g[i, j] = \frac{1}{W_{sb}} \sum_m \sum_n f[m, n] n_{\sigma_s}[i - m, j - n] n_{\sigma_b}(|f[m, n] - f[i, j]|)$$

Where:


$$n_{\sigma_s}[m, n] = \frac{1}{2\pi\sigma_s^2} e^{-\frac{1}{2}\left(\frac{m^2+n^2}{\sigma_s^2}\right)} \quad n_{\sigma_b}[k] = \frac{1}{2\pi\sigma_b^2} e^{-\frac{1}{2}\left(\frac{k^2}{\sigma_b^2}\right)}$$

$$W_{sb} = \sum_m \sum_n n_{\sigma_s}[i - m, j - n] n_{\sigma_b}(|f[m, n] - f[i, j]|)$$


Non-linear Operation

(Cannot be implemented using Convolution)

Gaussian vs. Bilateral Filtering: Example



Original




Gaussian
 $\sigma_s = 2$




Bilateral
 $\sigma_s = 2, \sigma_b = 10$

Gaussian vs. Bilateral Filtering: Example



Original




Gaussian
 $\sigma_s = 4$




Bilateral
 $\sigma_s = 4, \sigma_b = 10$

Gaussian vs. Bilateral Filtering: Example



Original



Gaussian
 $\sigma_s = 8$




Bilateral
 $\sigma_s = 8, \sigma_b = 10$

Bilateral Filtering: Changing σ_b




Bilateral

$$\sigma_s = 6, \sigma_b = 10$$



Bilateral

$$\sigma_s = 6, \sigma_b = 20$$



Bilateral

$$\sigma_s = 6, \sigma_b = \infty$$

(Gaussian Smoothing)

Template Matching



Template

How do we locate the template in the image?

Minimize:

$$E[i, j] = \sum_m \sum_n (f[m, n] - t[m - i, n - j])^2$$

$$E[i, j] = \sum_m \sum_n (f^2[m, n] + t^2[m - i, n - j] - 2f[m, n]t[m - i, n - j])$$

Maximize

Template Matching



Template

How do we locate the template in the image?

Maximize:

$$R_{tf}[i, j] = \sum_m \sum_n f[m, n] t[m - i, n - j] = t \otimes f$$

(Cross-Correlation)

Convolution vs. Correlation

Convolution:

$$g[i, j] = \sum_m \sum_n f[m, n] \underline{t[i - m, j - n]} = t * f$$


Correlation:

$$R_{tf}[i, j] = \sum_m \sum_n f[m, n] \underline{t[m - i, n - j]} = t \otimes f$$

No Flipping in Correlation

Problem with Cross-Correlation

$$R_{tf}[i, j] = \sum_m \sum_n f[m, n] t[m - i, n - j] = t \otimes f$$




$$R_{tf}(C) > R_{tf}(B) > R_{tf}(A)$$

We need $R_{tf}(A)$ to be the maximum!

Normalized Cross-Correlation

Account for energy differences

$$N_{tf}[i, j] = \frac{\sum_m \sum_n f[m, n] t[m - i, n - j]}{\sqrt{\sum_m \sum_n f^2[m, n]} \sqrt{\sum_m \sum_n t^2[m - i, n - j]}}$$





$$R_{tf}(A) > R_{tf}(B) > R_{tf}(C)$$

Normalized Cross-Correlation

Account for energy differences

$$N_{tf}[i, j] = \frac{\sum_m \sum_n f[m, n] t[m - i, n - j]}{\sqrt{\sum_m \sum_n f^2[m, n]} \sqrt{\sum_m \sum_n t^2[m - i, n - j]}}$$



Correlation: Issues

- Problem at borders
- Sensitive to object pose, scale and rotation
- Not good for general object recognition
- Good for feature detection
- Can be computationally expensive

References

Textbooks:

Computer Vision: Algorithms and Applications (Chapter 3.2)

Recommended Reading

Szelinski, 2011 (available online)

Digital Image Processing (Chapter 3)

González, R and Woods, R., Prentice Hall

Robot Vision (Chapter 6 and 7)

Horn, B. K. P., MIT Press

Computer Vision: A Modern Approach (Chapter 7)

Forsyth, D and Ponce, J., Prentice Hall

Papers:

[Tomasi 1998] C. Tomasi and R. Manduchi, "Bilateral filtering for gray and color images," in Proceedings of the IEEE International Conference on Computer Vision, 1998.